


Sarno lite · MPI + OpenMP + CUDA



low-gn · one exclusive node · one run per point

Native forcing; preparation excluded. GPU/CPU trajectories differ; see example discussion.